

Dual-Database Retrieval-Augmented Generation (RAG) Architecture & Workflow

Spring Boot 4.1.1 • Spring AI 2.0.1 • MongoDB 7.0 • Qdrant HNSW • Ollama Llama 3.3 (70B) & Mistral Large • SAP Fiori Horizon UI

📄 Document Ingestion Pipeline: From Upload to Indexing

1. PDF / Doc Upload

- Drag & drop file
- Auto-extract BibTeX
- Extract full text
- SHA-256 Checksum

DocumentRecord

2. Semantic Chunking

- Chunk size: ~500 chars
- 100 char sliding overlap
- Sentence boundary split
- Preserves context flows

Chunks: [C0, C1, ... Cn]

3. nomic-embed-text

- Local Ollama Embedding Model
- 768-dimensional float32
- Unit L2 normalized
- High-throughput batching

$v \in \mathbb{R}^{768} \ (||v|| = 1.0)$

🗄️ Metadata & Asset Storage (MongoDB)

- Stores full BSON, BibTeX, versions, PDF asset
- Caches dual Llama & Mistral summaries
- ACID Guarantees • B-Tree Index • Port 37017

📌 Qdrant Vector DB (Vector Index)

- Stores 768-D dense vectors + chunk payload
- Hierarchical Navigable Small World (HNSW)
- < 3ms Cosine Distance ANN • Port 16333/16334

🔧 Concurrent Summarization

1. Llama 3.3 (70B Instruct):
Analytical synthesis • Methodologies
2. Mistral Large (2411):
Executive brief • Key findings

Concurrent Async Threads
Duration tracked & rendered in Object Page

🗨️ Query & Retrieval Workflow: From Question to Answer

📄 Input: Natural Language Question

Natural Language Question:

"What are the primary findings of this paper?"

- Emitted from Object Page
- POST /api/v1/documents/:id/chat
- Passes prior chat history

RestBackendAdapter

Keycloak Bearer JWT Header
nginx Ingress → Spring Boot

🔍 Processing: Query Vectorization & Filtering

1. Query Vectorization:

$v_{\text{query}} = \text{Embed}(\text{question})$

2. Filtered Cosine Distance:

$\text{doc_id} = \text{currentDocument}$
Threshold: $\cos(\theta) > 0.65$

3. Top-K Extraction (K=3):

- Chunk #2 (Score: 0.94) • Page 4
"Multi-head attention allows..."
- Chunk #0 (Score: 0.88) • Page 1
"We propose the Transformer..."
- Chunk #7 (Score: 0.73) • Page 9

🗨️ Context: Retrieved Passages & Chat History

[SYSTEM PROMPT]

You are Spring AI Assistant powered by Llama 3.3. Ground responses strictly in the retrieved passages with citations.

[RETRIEVED PASSAGES]

[Excerpt 1] Multi-Head Attention ...
[Excerpt 2] We propose the Transformer ...

[RECENT CHAT TURNS]

User: Tell me about parallelization.
Assistant: Unlike RNNs, Transformer ...

[USER QUESTION]

"What are the primary findings?"

Zero-Hallucination Guardrail Active

🗨️ Output: Streaming Synthesized Response

Streaming Synthesized Response (Markdown):

"The primary findings demonstrate that the **Transformer** replaces recurrent layers with **Multi-Head Attention** [Excerpt 1]. On WMT 2014 translation benchmarks, the big model achieves 28.4 BLEU, training in 3.5 days on 8 GPUs [Excerpt 2]."

✓ Model: Llama 3.3 (70B Instruct) via Spring AI • Latency: 1.8s
Directly verifiable in SAP UI5 Object Page with interactive excerpt cards

📄 Audit Trail: Relevancy & Source Tracking

- Chunk 2 (Relevance: 94.2%) – "Multi-Head Attention ..."
 - Chunk 0 (Relevance: 88.0%) – "Attention Is All You Need ..."
- Full audit trail preserved in MongoDB document version logs